

Global Stock Market Analytics

Jerry Kiely

In this project we will look at predicting the open direction of the Nifty 50 index by looking at other indices and indicators. We will break the project up into five parts:

- ① preparing the master data from the global indices
- ② preliminary analysis of the data
- ③ predictive modelling of open direction of Nifty 50
- ④ comparing different models at predicting open direction
- ⑤ sentiment analysis of X / Twitter data relating to Nifty 50

Phase 1 - Prepare the Master Data

The indexes of interest are:

- NSEI: Nifty 50
- DJI: Dow Jones Index
- IXIC: Nasdaq
- HSI: Hang Seng
- N225: Nikkei 225
- GDAXI: Dax
- VIX: Volatility Index

Phase 1 - Prepare the Master Data

Global
Stock
Market
Analytics

Jerry Kiely

We download and merge the required data, using LOCF to impute missing data, and adding variables for MONTH, QUARTER, and YEAR.

```
def retrieve_data(index, start_date = '2017-12-1', end_date = '2024-1-31', progress = False):
    data = yf.download(f'{index}', start_date, end_date, progress = progress)

    # create daily returns for each index
    data['Daily Returns'] = data.Close.pct_change() * 100

    # rename columns - prefix with index name
    data.columns = ["_".join(c.upper() for c in column.split()) for column in data.columns]
    data.columns = [f"{index}_{column}" for column in data.columns]

    return data

data = [retrieve_data(index) for index in INDICES]

# merge data with outer join
merged = pd.concat(data, axis = 1)

# impute missing data using LOCF (forward fill)
merged.ffmpeg(inplace = True)

# add indicators for MONTH, QUARTER, and YEAR
merged['MONTH'] = merged.index.month
merged['QUARTER'] = merged.index.quarter
merged['YEAR'] = merged.index.year
```

Phase 2 - Preliminary Analysis






We perform some preliminary analysis on the master data hoping to answer the following questions:

- ① Which index has given consistently good returns?
- ② Which index was highly volatile?
- ③ How are global markets correlated during 6 years period and is the correlation structure similar in the most recent year - i.e. 2023?
- ④ Assuming a primary target variable of “Nifty Opening Price Direction”, what are preliminary insights?

Phase 2 - Preliminary Analysis







Looking at the Q-Q Plots, the daily returns do not appear to follow - or be drawn from - a Normal Distribution - specifically at the tails. But normality is not a requisite or assumption of Logistic Regression with respect to independent variables.

Q-Q Plots of Daily Returns



Phase 2 - Which index has given consistently good returns?

Box Plots grouped by Year



All indexes seem pretty consistent - all years have similar spreads, and consistent medians, with one or two exceptions. All indexes for 2020 have more outliers than normal. But HSI seems to have more outliers in 2022 than in 2020.

Phase 2 - Which index has given consistently good returns?

Looking at the summary statistics, the clear winner here is NSEI - not one year in the range has a negative mean return. And with the exception of 2020 and 2022, NSEI has low volatility (< 1) throughout all years.

NSEI

	count	mean	std	var
YEAR				
2018	260	0.012	0.804	0.647
2019	260	0.062	0.862	0.744
2020	262	0.059	2.004	4.015
2021	261	0.094	0.980	0.960
2022	260	0.055	1.096	1.202
2023	260	0.079	0.620	0.384

Phase 2 - Which index was highly volatile?

As for the most volatile indexes, it's a toss between IXIC and HSI, both of whom have high volatility (> 1) when compared to the other indexes.

IXIC

	count	mean	std	var
YEAR				
2018	260	-0.020	1.330	1.768
2019	260	0.133	0.975	0.950
2020	262	0.170	2.200	4.838
2021	261	0.096	1.124	1.262
2022	260	-0.124	2.000	4.001
2023	260	0.157	1.085	1.177



HSI

	count	mean	std	var
YEAR				
2018	260	-0.035	1.244	1.547
2019	260	0.033	0.981	0.962
2020	262	0.026	1.445	2.087
2021	261	-0.028	1.262	1.593
2022	260	-0.021	2.054	4.221
2023	260	-0.053	1.409	1.984

Phase 2 - Which index has given consistently good returns?

Looking at bar plots for median returns by year, again the clear winner here is NSEI - at no time is the median daily returns for any of the years below 0. IXIC has an unusually high 2020, but a bad 2022. HSI also has an unusually bad 2022.

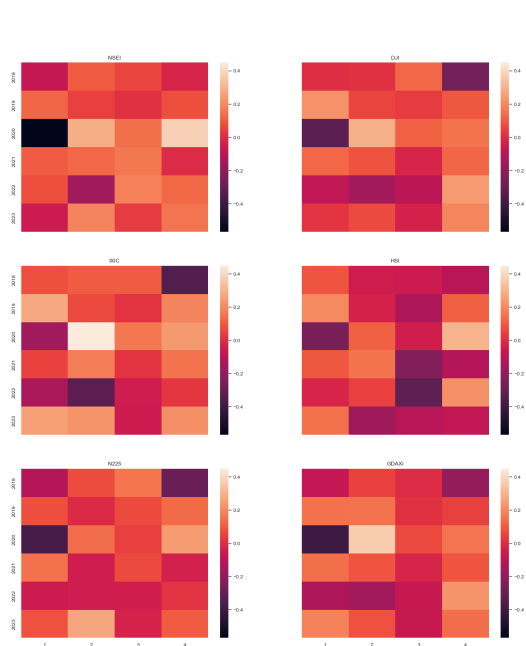
Bar Plots of Median Returns grouped by Year



Phase 2 - Which index has given consistently good returns?

Looking at heat maps of mean returns, with the exception of the 1st quarter in 2020, NSEI has pretty consistent daily returns - where most cells are pretty bright, denoting above 0. Most of the other indexes have a blend of light and dark, which would indicate more volatile behaviour over the quarters.

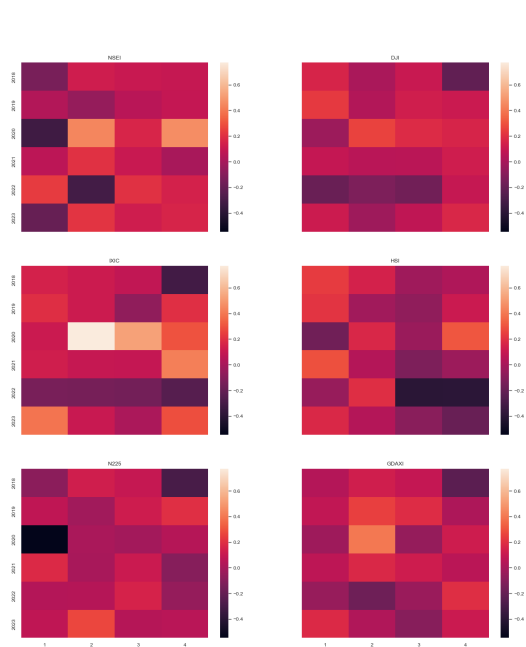
Heat Maps of Mean Returns grouped by Year



Phase 2 - Which index has given consistently good returns?


On the other hand, when looking at median returns across quarters NSEI seems pretty average - there does not seem to be a clear winner here.

Heat Maps of Median Returns grouped by Year



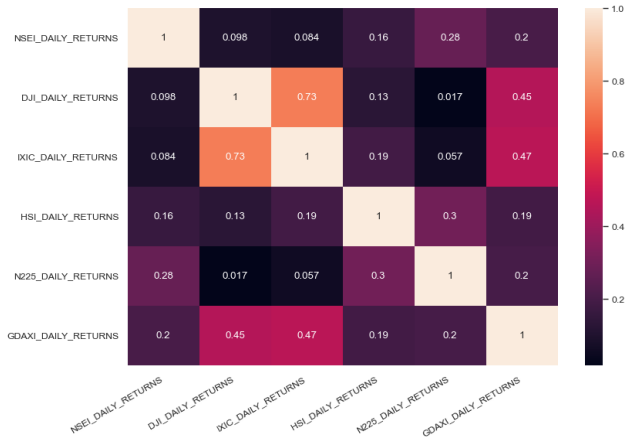
Phase 2 - How are global markets correlated during the 6 years period?

It looks like strong correlation between daily returns of IXIC and DJI, and some correlation between GDAXI and DJI. These indexes are likely to result in multicollinearity at the regression stage.



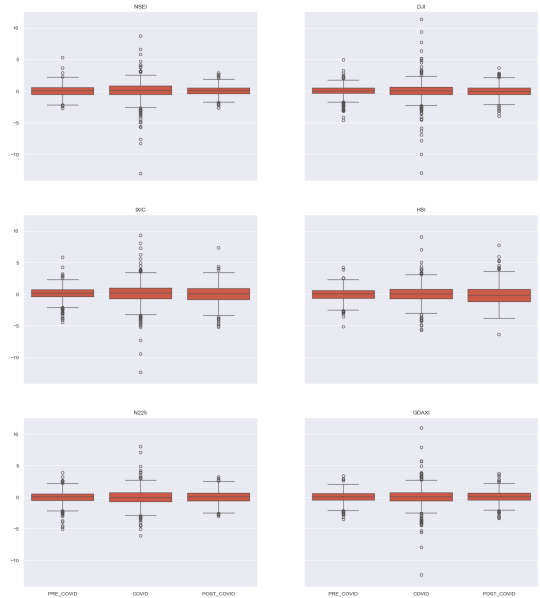
Phase 2 - Is the correlation structure similar in the recent year - 2023?

We can see similar - but slightly weaker - correlations exist between the same indexes for 2023.



Phase 2 - Covid Period

Box Plots grouped by Pandemic Period



We can see that the spreads of each index over the Pandemic are consistent, with the Covid period itself having more outliers - which of course you might expect.

All indexes had higher volatility over the Covid period.


- NSEI performed reasonably well over the Covid period, with an increase in volatility in the period, and with a significant bump in the Post Covid period.
- DJI seemed consistent over the three periods, with an increase in volatility in the Covid period.
- IXIC looked pretty good over the three period, but maybe slightly more volatile overall, and in particular in the Covid period.
- HSI has performed poorly in general, with negative returns in the pre and post Covid periods, and with consistently greater volatility than most.
- N225 appears to perform not so well, and with relatively high volatility.
- GDAXI also appears to perform not so well in general, and with relatively high volatility.

Phase 2 - Covid Period

Bar Plots grouped by Pandemic Period

20

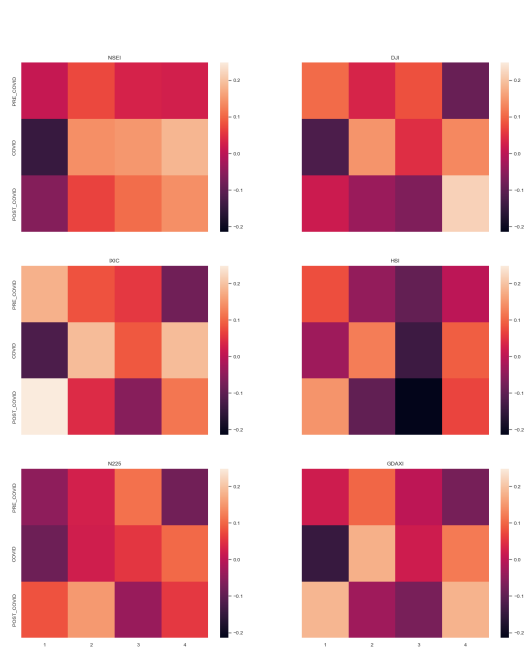
With respect to returns, we can see that IXIC looks like the clear winner, with NSEI in second place, and DJI and GSAXI in a fight for third place. HSI appears to have had a terrible Post Covid period, and N225 appears to have had a pretty bad Covid period.



Phase 2 - Covid Period

Again, NSEI appears to be the most consistent of all indexes. All indexes have bad first quarters during the Pandemic, but improve post Covid.


Heat Maps of Mean Returns grouped by Pandemic Period



Phase 2 - Covid Period

Median returns tells a similar story over the Pandemic period - HSI in particular appears to have had the worst recovery.

Heat Maps of Median Returns grouped by Pandemic Period



We try to estimate the time taken for each of the indexes to return to the Pre Covid levels - the approach is to find how many days it takes for each index to reach a value greater than or equal to the Pre Covid mean returns value.

NSEI returned to pre-covid levels on 2022-05-16 after 7 trading day(s)

DJI returned to pre-covid levels on 2022-05-13 after 6 trading day(s)

IXIC returned to pre-covid levels on 2022-05-10 after 3 trading day(s)

HSI returned to pre-covid levels on 2022-05-11 after 4 trading day(s)

N225 returned to pre-covid levels on 2022-05-06 after 1 trading day(s)

GDAXI returned to pre-covid levels on 2022-05-10 after 3 trading day(s)

Interestingly, N225 returned to it's Pre Covid level after just 1 day.

Phase 2 - Nifty Opening Price Direction: Preliminary Insights?

We define the Nifty Opening Price Direction - NSEI_OPEN_DIR - as 1 if the value of the NSEI Opening price at time t is greater than the value of the NSEI Closing price at time $t - 1$, and 0 otherwise. Lets look at the percentages of $\text{NSEI_OPEN_DIR} = 1$ by year:

Nifty Fifty Daily Movement

YEAR

2018	70.38%
2019	69.23%
2020	70.61%
2021	71.65%
2022	59.23%
2023	67.31%

With the exception of 2022, every year has between 67% and 72% where $\text{NSEI_OPEN_DIR} = 1$.

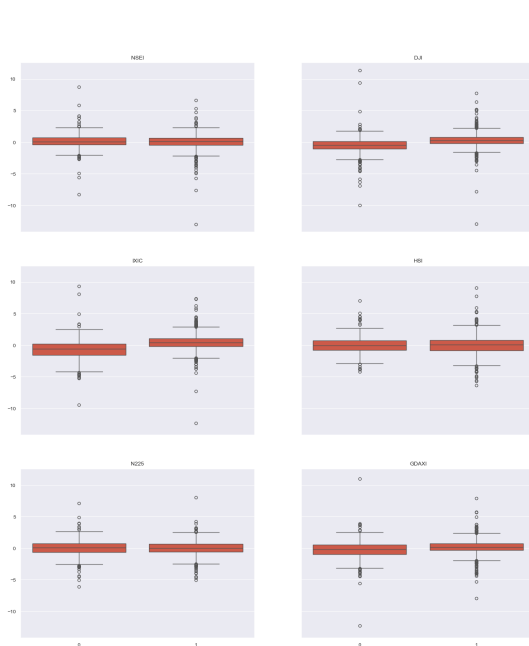
Phase 2 - Nifty Opening Price Direction

Global
Stock
Market
Analytics

Jerry Kiely

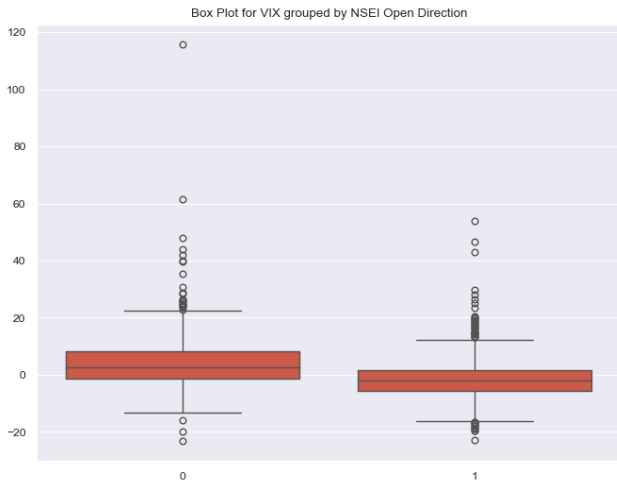
Looking at the the box plots, the data looks fairly consistent across each category of NSEI_OPEN_DIR, with the exceptions of DJI, IXIC and VIX.

Box Plots grouped by NSEI Open Direction



Phase 2 - Nifty Opening Price Direction

We look at VIX separately - as it requires a different scale.



Phase 3 - Logistic Model

Before proceeding with modelling NSEI_OPEN_DIR, lets define, and add, some ratios and indicators, making use of the Python technical analysis library:

```
master["NSEI_HL_RATIO"] = master["NSEI_HIGH"] / master["NSEI_LOW"]
master["DJI_HL_RATIO"]  = master["DJI_HIGH"]  / master["DJI_LOW"]

master["NSEI_RSI"]      = ta.momentum.rsi(master["NSEI_CLOSE"])
master["DJI_RSI"]       = ta.momentum.rsi(master["DJI_CLOSE"])

master["NSEI_TSI"]      = ta.momentum.tsi(master["NSEI_CLOSE"])
master["DJI_TSI"]       = ta.momentum.tsi(master["DJI_CLOSE"])
```


Phase 3 - Logistic Model

Global
Stock
Market
Analytics

Jerry Kiely

We define a function that will prune any features that are either found to be insignificant, or that are found to be collinear:

```
def prune(X, y, verbose = True):
    dropped = []
    while True:
        model = Logit(y, X).fit(displ = 0)
        exog = model.model.exog
        names = model.model.exog_names

        insignificant = [p for p in zip(model.pvalues.index[1:], model.pvalues[1:]) if p[1] > 0.05]

        values = [variance_inflation_factor(exog, i) for i in range(1, exog.shape[1])]
        colinear = [val for val in zip(names[1:], values) if val[1] > 5]

        if insignificant:
            insignificant.sort(key = lambda p: -p[1])

            if verbose:
                print(f"dropping {insignificant[0][0]} with p-value {insignificant[0][1]}")

            X = X.drop([insignificant[0][0]], axis = 1)
            dropped.append(insignificant[0][0])

        elif colinear:
            colinear.sort(key = lambda c: -c[1])

            if verbose:
                print(f"dropping {colinear[0][0]} with vif {colinear[0][1]}")

            X = X.drop([colinear[0][0]], axis = 1)
            dropped.append(colinear[0][0])

        else:
            return model, dropped

model, dropped = prune(X_train, y_train)
```

The function outputs a list of pruned features, together with the associated p-value or vif value. The function returns the pruned model, together with a list of pruned feature names.

```
dropping    DJI_DAILY_RETURNS with p-value 0.7234766099770011
dropping GDAXI_DAILY_RETURNS with p-value 0.6162105670376612
dropping      NSEI_HL_RATIO with p-value 0.4277618505298021
dropping      DJI_HL_RATIO with p-value 0.1563055988923202
dropping NSEI_DAILY_RETURNS with p-value 0.13281329048460666
dropping      NSEI_TSI with      vif 5.865700460659149
dropping      NSEI_RSI with p-value 0.7783762272653001
```

Phase 3 - Logistic Model

The Logistic Model Summary

Dep. Variable:	NSEI_OPEN_DIR	No. Observations:	1220
Model:	Logit	Df Residuals:	1213
Method:	MLE	Df Model:	6
Date:	Wed, 10 Jul 2024	Pseudo R-squ.:	0.1375
Time:	10:31:02	Log-Likelihood:	-660.02
converged:	True	LL-Null:	-765.23
Covariance Type:	nonrobust	LLR p-value:	1.141e-42

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-1.4041	0.656	-2.139	0.032	-2.690	-0.118
IXIC_DAILY_RETURNS	0.4552	0.075	6.093	0.000	0.309	0.602
HSI_DAILY_RETURNS	-0.1395	0.053	-2.632	0.008	-0.243	-0.036
N225_DAILY_RETURNS	-0.1960	0.068	-2.897	0.004	-0.329	-0.063
VIX_DAILY_RETURNS	-0.0397	0.013	-3.054	0.002	-0.065	-0.014
DJI_RSI	0.0447	0.013	3.415	0.001	0.019	0.070
DJI_TSI	-0.0205	0.008	-2.660	0.008	-0.036	-0.005

The final model is:



$$\ln\left(\frac{p}{1-p}\right) = -1.4041 + 0.4552x_1 - 0.1395x_2 - 0.1960x_3 - 0.0397x_4 + 0.0447x_5 - 0.0205x_6$$

where:

variable		value
x_1	IXIC_DAILY_RETURNS	
x_2	HSI_DAILY_RETURNS	
x_3	N225_DAILY_RETURNS	
x_4	VIX_DAILY_RETURNS	
x_5	DJI_RSI	
x_6	DJI_TSI	

Phase 3 - Logistic Model

We plot the ROC curve for the train and test data:



Phase 3 - Logistic Model

The classification report and confusion matrix for the train data:

Train Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.53	0.68	0.60	391
1.0	0.83	0.72	0.77	829
accuracy			0.70	1220
macro avg	0.68	0.70	0.68	1220
weighted avg	0.73	0.70	0.71	1220

Train Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	265	234
1	126	595

Phase 3 - Logistic Model

The classification report and confusion matrix for the test data:

Test Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.53	0.65	0.58	97
1.0	0.82	0.73	0.77	208
accuracy			0.70	305
macro avg	0.67	0.69	0.67	305
weighted avg	0.72	0.70	0.71	305

Test Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	63	57
1	34	151

We compare some statistics across train and test data:

Train Data - Sensitivity: 71.77%

Test Data - Sensitivity: 72.6%

Train Data - Specificity: 67.77%

Test Data - Specificity: 64.95%

Train Data - AUC ROC: 0.753

Test Data - AUC ROC: 0.752

The AUC is consistent across the train and test data, and the sensitivity and specificity values are reasonably good, and somewhat consistent across the train and test data. However, the accuracy of the model is not great. We could potentially obtain better results by selecting a different approach to classification.

We compare the performance of a number of different models to see if we can improve on the accuracy of our original model:

- Logistic Regression
- Naive Bayes
- KNN
- Decision Tree
- Random Forest
- SVM
- MLP
- Deep Learning (PyTorch)



Phase 4 - Compare Models

Global
Stock
Market
Analytics

Jerry Kiely

Logistic Regression

ROC Curves



Phase 4 - Compare Models

Logistic Regression

Train Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.58	0.61	0.59	391
1.0	0.81	0.79	0.80	829
accuracy			0.73	1220
macro avg	0.70	0.70	0.70	1220
weighted avg	0.74	0.73	0.73	1220

Train Data - Confusion Matrix:

col_0	0.0	1.0
row_0		
0	239	174
1	152	655

Phase 4 - Compare Models

Logistic Regression

Test Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.56	0.63	0.59	97
1.0	0.82	0.77	0.79	208
accuracy			0.72	305
macro avg	0.69	0.70	0.69	305
weighted avg	0.73	0.72	0.73	305

Test Data - Confusion Matrix:

col_0	0.0	1.0
row_0		
0	61	48
1	36	160

Logistic Regression

Train Data - Sensitivity: 79.01%

Test Data - Sensitivity: 76.92%

Train Data - Specificity: 61.13%

Test Data - Specificity: 62.89%

Train Data - AUC ROC: 0.758

Test Data - AUC ROC: 0.767



We notice that the scikit-learn Logistic Regression model slightly outperforms the Statsmodels Logit model.

We use the Statsmodels Logit model when we need to perform analysis of the model and features, then move to scikit-learn's LogisticRegression model after the model has been finalised.

Phase 4 - Compare Models

Naive Bayes

ROC Curves



Phase 4 - Compare Models

Naive Bayes

Train Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.50	0.68	0.57	391
1.0	0.82	0.68	0.74	829
accuracy			0.68	1220
macro avg	0.66	0.68	0.66	1220
weighted avg	0.71	0.68	0.69	1220

Train Data - Confusion Matrix:

col_0	0.0	1.0
row_0		
0	264	266
1	127	563

Phase 4 - Compare Models

Naive Bayes

Test Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.47	0.70	0.56	97
1.0	0.82	0.62	0.71	208
accuracy			0.65	305
macro avg	0.64	0.66	0.63	305
weighted avg	0.71	0.65	0.66	305

Test Data - Confusion Matrix:

col_0	0.0	1.0
row_0		
0	68	78
1	29	130

Naive Bayes

Train Data - Sensitivity: 67.91%

Test Data - Sensitivity: 62.5%

Train Data - Specificity: 67.52%

Test Data - Specificity: 70.1%

Train Data - AUC ROC: 0.728

Test Data - AUC ROC: 0.702



Phase 4 - Compare Models

Global
Stock
Market
Analytics

Jerry Kiely

KNN

ROC Curves



Phase 4 - Compare Models

KNN

Train Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.47	0.74	0.57	391
1.0	0.83	0.60	0.70	829
accuracy			0.65	1220
macro avg	0.65	0.67	0.64	1220
weighted avg	0.71	0.65	0.66	1220

Train Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	288	328
1	103	501

Phase 4 - Compare Models

KNN

Test Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.47	0.69	0.56	97
1.0	0.81	0.63	0.71	208
accuracy			0.65	305
macro avg	0.64	0.66	0.63	305
weighted avg	0.70	0.65	0.66	305

Test Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	67	77
1	30	131

KNN

Train Data - Sensitivity: 60.43%

Test Data - Sensitivity: 62.98%

Train Data - Specificity: 73.66%

Test Data - Specificity: 69.07%

Train Data - AUC ROC: 0.745

Test Data - AUC ROC: 0.716



Phase 4 - Compare Models

Global
Stock
Market
Analytics

Jerry Kiely

Decision Tree

ROC Curves



Phase 4 - Compare Models

Decision Tree

Train Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.41	0.86	0.56	391
1.0	0.87	0.41	0.56	829
accuracy			0.56	1220
macro avg	0.64	0.64	0.56	1220
weighted avg	0.72	0.56	0.56	1220

Train Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	338	488
1	53	341

Phase 4 - Compare Models

Decision Tree

Test Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.40	0.84	0.54	97
1.0	0.84	0.42	0.56	208
accuracy			0.55	305
macro avg	0.62	0.63	0.55	305
weighted avg	0.70	0.55	0.55	305

Test Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	81	121
1	16	87

Decision Tree

Train Data - Sensitivity: 41.13%

Test Data - Sensitivity: 41.83%

Train Data - Specificity: 86.45%

Test Data - Specificity: 83.51%



Train Data - AUC ROC: 0.719

Test Data - AUC ROC: 0.712

Phase 4 - Compare Models

Random Forest

ROC Curves



Phase 4 - Compare Models

Random Forest

Train Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.54	0.76	0.63	391
1.0	0.86	0.69	0.77	829
accuracy			0.71	1220
macro avg	0.70	0.73	0.70	1220
weighted avg	0.76	0.71	0.72	1220

Train Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	297	254
1	94	575

Phase 4 - Compare Models

Random Forest

Test Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.51	0.72	0.60	97
1.0	0.84	0.68	0.75	208
accuracy			0.69	305
macro avg	0.68	0.70	0.67	305
weighted avg	0.73	0.69	0.70	305

Test Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	70	67
1	27	141

Phase 4 - Compare Models

Random Forest

Train Data - Sensitivity: 69.36%

Test Data - Sensitivity: 67.79%

Train Data - Specificity: 75.96%

Test Data - Specificity: 72.16%

Train Data - AUC ROC: 0.796

Test Data - AUC ROC: 0.756



Phase 4 - Compare Models

Global
Stock
Market
Analytics

Jerry Kiely

SVM

ROC Curves



Phase 4 - Compare Models

SVM

Train Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.62	0.57	0.59	391
1.0	0.80	0.84	0.82	829
accuracy			0.75	1220
macro avg	0.71	0.70	0.71	1220
weighted avg	0.75	0.75	0.75	1220

Train Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	222	136
1	169	693

Phase 4 - Compare Models

SVM

Test Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.61	0.59	0.60	97
1.0	0.81	0.82	0.82	208
accuracy			0.75	305
macro avg	0.71	0.70	0.71	305
weighted avg	0.75	0.75	0.75	305

Test Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	57	37
1	40	171

SVM

Train Data - Sensitivity: 83.59%

Test Data - Sensitivity: 82.21%

Train Data - Specificity: 56.78%

Test Data - Specificity: 58.76%

Train Data - AUC ROC: 0.760

Test Data - AUC ROC: 0.763



Phase 4 - Compare Models

Global
Stock
Market
Analytics

Jerry Kiely

MLP

ROC Curves



Phase 4 - Compare Models

MLP

Train Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.57	0.68	0.62	391
1.0	0.83	0.76	0.80	829
accuracy			0.73	1220
macro avg	0.70	0.72	0.71	1220
weighted avg	0.75	0.73	0.74	1220

Train Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	264	198
1	127	631

Phase 4 - Compare Models

MLP

Test Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.55	0.67	0.60	97
1.0	0.83	0.75	0.78	208
accuracy			0.72	305
macro avg	0.69	0.71	0.69	305
weighted avg	0.74	0.72	0.73	305

Test Data - Confusion Matrix:

NSEI_OPEN_DIR	0.0	1.0
row_0		
0	65	53
1	32	155

Phase 4 - Compare Models

MLP

Train Data - Sensitivity: 76.12%

Test Data - Sensitivity: 74.52%

Train Data - Specificity: 67.52%



Test Data - Specificity: 67.01%

Train Data - AUC ROC: 0.773

Test Data - AUC ROC: 0.757

Phase 4 - Compare Models

Deep Learning (PyTorch)



Phase 4 - Compare Models

Deep Learning (PyTorch)

Train Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.71	0.56	0.62	391
1.0	0.81	0.89	0.85	829
accuracy			0.79	1220
macro avg	0.76	0.73	0.74	1220
weighted avg	0.78	0.79	0.78	1220

Train Data - Confusion Matrix:

col_0	0.0	1.0
row_0		
0	218	89
1	173	740

Phase 4 - Compare Models

Deep Learning (PyTorch)

Test Data - Classification Report:

	precision	recall	f1-score	support
0.0	0.66	0.55	0.60	97
1.0	0.80	0.87	0.84	208
accuracy			0.77	305
macro avg	0.73	0.71	0.72	305
weighted avg	0.76	0.77	0.76	305

Test Data - Confusion Matrix:

col_0	0.0	1.0
row_0		
0	53	27
1	44	181

Deep Learning (PyTorch)

Train Data - Sensitivity: 87.02%

Test Data - Sensitivity: 89.26%

Train Data - Specificity: 54.64%

Test Data - Specificity: 55.75%

Train Data - AUC ROC: 0.772

Test Data - AUC ROC: 0.738

Phase 4 - Compare Models

```
Naive Bayes - Test Data - AUC ROC: 0.702
Decision Tree - Test Data - AUC ROC: 0.712
      KNN - Test Data - AUC ROC: 0.716
Deep Learning - Test Data - AUC ROC: 0.738
Random Forest - Test Data - AUC ROC: 0.756
      MLP - Test Data - AUC ROC: 0.757
      SVM - Test Data - AUC ROC: 0.763
Logistic Regression - Test Data - AUC ROC: 0.767
```

Interestingly, the values of AUC for each model are not indicative of the accuracy of the models - the Logistic Regression model, which had a very average accuracy, has the highest AUC, while the Deep Learning model, the model with the highest accuracy, has an average AUC.

Phase 5 - Sentiment Analysis

We now turn to Twitter / X data relating to the Nifty 50 index to see if we can mine some sentiment.

We load the tweets, create a data frame, and then do some basic pre-processing of the data to:

- 1 transform all words to lowercase
- 2 remove all punctuation
- 3 remove all digits
- 4 remove stopwords

We load the tweets:

```
with open(os.path.join(os.getcwd(), "Tweets.txt")) as file:
    tweets = [line.rstrip() for line in file]

data = pd.DataFrame(
    [line for line in tweets if len(line) > 0],
    columns= ["Tweets"]
)
data.head()
```

Tweets

```
0 #bankNifty 50100 ce looks good at 70+-2 for a ...
1 "#market #banknifty #OptionsTrading #optionbuy...
2 PENNY STOCK MADHUCON PROJECTS LTD cmp-11 FOLLO...
3 #Nifty50 has been in a healthy uptrend since t...
4 #Gravita #livetrading #stockstowatch #stocksin...
```

Phase 5 - Sentiment Analysis

We perform the basic transformation:

```
stop_words = set(stopwords.words('english'))
remove_punc = str.maketrans('', '', punctuation)
remove_digits = str.maketrans('', '', digits)

def preprocess_tweet(tweet):
    tokens = word_tokenize(
        tweet.lower().translate(remove_punc).translate(remove_digits)
    )
    return " ".join([word for word in tokens if word not in stop_words])

cleaned = data["Tweets"].apply(preprocess_tweet)
cleaned.head()
```

```
0          banknifty ce looks good target nifty nifty
1    market banknifty optionstrading optionbuying t...
2    penny stock madhucon projects ltd cmp followht...
3    nifty healthy uptrend since beginning year did...
4    gravita livetrading stockstowatch stocksinfocu...
```

Phase 5 - Sentiment Analysis

We look at the top 10 words by frequency:

```
tweet_words = cleaned.str.cat(sep = " ")

freq_dist = nltk.FreqDist(tweet_words.split())
word_freq = pd.DataFrame(freq_dist.most_common(10), columns=["Word", "Freq"])
word_freq.head(10)
```

	Word	Freq
0	nifty	399
1	banknifty	104
2	stockmarket	71
3	niftybank	45
4	stockmarketindia	44
5	sensex	43
6	stocks	38
7	optionstrading	36
8	bse	34
9	breakoutstocks	31
...		

Phase 5 - Sentiment Analysis

We include the word “nifty” to list of words to remove from the tweets:

```
stop_words = set(stopwords.words('english')) | set(["nifty"])

data["Cleaned_Tweets"] = data["Tweets"].apply(preprocess_tweet)
data.head()
```

```

                                Tweets \
0  #bankNifty 50100 ce looks good at 70+-2 for a ...
1  "#market #banknifty #OptionsTrading #optionbuy...
2  PENNY STOCK MADHUCON PROJECTS LTD cmp-11 FOLLO...
3  #Nifty50 has been in a healthy uptrend since t...
4  #Gravita #livetrading #stockstowatch #stocksin...
```

```

                                Cleaned_Tweets
0                                banknifty ce looks good target
1  market banknifty optionstrading optionbuying t...
2  penny stock madhucon projects ltd cmp followht...
3  healthy uptrend since beginning year didnt bre...
4  gravita livetrading stockstowatch stocksinfocu...
```


We should review more of the words as candidates to include for removal.

Phase 5 - Sentiment Analysis

Global
Stock
Market
Analytics

Jerry Kiely

We visualise these top 20 words by frequency:



Phase 5 - Sentiment Analysis

We extract sentiment scores for the tweets:

```
sia = SentimentIntensityAnalyzer()

scores = data["Cleaned_Tweets"].apply(lambda x: sia.polarity_scores(x))

data["Positive_Score"] = scores.apply(lambda x: x["pos"])
data["Negative_Score"] = scores.apply(lambda x: x["neg"])
data["Neutral_Score"] = scores.apply(lambda x: x["neu"])
data["Compound_Score"] = scores.apply(lambda x: x["compound"])
```

Phase 5 - Sentiment Analysis

We look at the summary statistics for the sentiment scores:

	Positive_Score	Negative_Score	Neutral_Score	Compound_Score
count	245.000000	245.000000	245.000000	245.000000
mean	0.116216	0.028490	0.855314	0.172913
std	0.146029	0.071052	0.156892	0.343955
min	0.000000	0.000000	0.213000	-0.807400
25%	0.000000	0.000000	0.742000	0.000000
50%	0.053000	0.000000	0.868000	0.000000
75%	0.194000	0.000000	1.000000	0.440400
max	0.787000	0.405000	1.000000	0.928700


The compound score has a reasonable spread with a median of 0.

Phase 5 - Sentiment Analysis

Global
Stock
Market
Analytics

Jerry Kiely

We plot the sentiment scores:



Phase 5 - Sentiment Analysis

It looks like the compound sentiment scores, if slightly positively skewed, could be useful as an extra feature. This positive skewing could be understood in terms of the distribution of open directions in the data.

But in the end, without access to historical tweets, it would be impossible to judge the efficacy of the sentiment scores conclusively. Nevertheless it would appear promising, and it should be investigated further.

Predicting NSEI open direction is a very interesting problem. We have shown the ability to train models with accuracy of around 75% in the case of the SVM model, and 77% - 79% in the case of the Deep Learning (PyTorch) model - I am confident that with enhanced model tuning, and by adding extra features such as sentiment compound score, we could improve the model accuracy significantly.